

CS50300: Operating Systems

LAB1 ANSWERS

Zichen Wang
wang4113@purdue.edu

September 5, 2018

1 Objectives

2 Readings

3 Running XINU

3.1

An alternative way to ‘idle’ is to call **halt()** defined in **intr.S**. The assembly code of ‘halt’ only jumps to itself, which means looping forever.

3.2

The **welcome()** function is under **system/welcome.c**.

4 XINU system calls

4.1

The C type of syscall, defined in **include/kernel.h**, is **int**.

4.2

If the parent process has terminated, the process table field **prstate** should be *PR_FREE*.

Question: What should we do if the parent process has terminated before the child process?

In x86, when a process terminates, its child process will change the parent to ‘init’ process. However, I cannot find such operation in system/kill.c.

5 XINU's run-time environment

5.1 Changing byte order using assembly code

The instruction **bswapl** can change byte order of a 32-bit register directly.

I rename the file and function name of the C function version to **revbyteorder_cfun**, in order to avoid conflict with **revbyteorder.S** when testing.

5.2 Checking segment boundaries

The addresses of the end of the text, data, and bss segments of the XINU OS are defined by $\&etext - 1$, $\&edata - 1$ and $\&ebss - 1$ respectively.

Hints from initialize.c.

5.3 Run-time stack: process creation and function call

We will use the inline assembly code to print the top of run-time stack, by copying the value of the register **%esp** and the content inside (**%esp**) into two variables.

The address and content of the top of the run-time stack inside **main()** will remain the **same** at the time before **myprogA()** is created and after **myprogA()** has been created and resumed. It is because **myprocA()** is a new independent process, and able to run along with **main()** process. The stack of **myprocA()** will be created in the free memory area when the process spawns and will **not** change the stack structure of its parent process. Therefore, in these two printing, **main()** process's stack and **main()** function's stack frame are being printed, and the results are same.

When **myprogA()** calls the function **myfuncA()**, things will be different. Before **myfuncA()** is called, **myprocA()** process's stack and **myprocA()** function's stack frame are being printed. Then, **myprogA()** will use its own stack to maintain a function call to **myfuncA()**, pushing the arguments of **myfuncA()** and return address to the stack. After that, **myprogA()** will jump to **myfuncA()**, but the process remains unchanged. In function **myfuncA()**, the stack belongs to **myprogA()** process. So, after **myfuncA()** is called, **myprocA()** process's stack and **myfuncA()** function's stack frame are being printed.

5.4 Comparing two run-time stacks

6 Hijacking a process via stack smashing

The strategy for finding the return address of **sleepms()** function is very simple. In **myfuncA()** process, it is easy to get the process id of **myprogA()** process, i.e. PPID of **myfuncA()** process. Then, we use the PPID to retrieve the stack parameters of **myprogA()** process and the whole stack from **prstkbase** to **prstkptr** can be printed. There are too many addresses at the first glance on the stack information, but the system call **sleepms** has a parameter '3000'. From the system perspective, when a process make a system call or function call, it will push the return address just following the arguments of the function. According to this convention, I find that '0x00000BB8'

(‘3000’ in hex) in the stack and the content under ‘0x0000BB8’ is just the return address of *sleepms* function. In my version, the return address is ‘0x00103491’, and the location of the return address is *prstkbase* − 52.

However, things will not finish. In order to let *myprocA()* process exit normally, i.e. its resources can be released and other process with lower priority can continue to run, *myprocA()* process should return to *INITRET* address after running the malware code. From the convention, the process will use **ret** instruction to pop the return address to **eip** register. Therefore, we may put the *INITRET* address before the malware code address in the stack. The *INITRET* address is located in *prstkbase* − 4 defined by *create.c*, so in my version $*(int*)(prstkbase-48) = *(int*)(prstkbase-4)$ will work.

7 Bonus problem

As detailed above, we may put any address before the malware code address in the stack to control *myprocA()* process’s **eip** register when returning from the malware code. In my version, to minimize the disruption on *myprocA()* process, we should move the normal address from *prstkbase* − 52 to *prstkbase* − 48, and then put the malware address to *prstkbase* − 52. By doing so, the context will return to *myprocA()* process normally after running the malware code, and continue to print its stack parameters.