

CS50300: Operating Systems

LAB1 ANSWERS

Zichen Wang
wang4113@purdue.edu

September 12, 2018

1 Objectives

2 Readings

3 Running XINU

3.1

An alternative way to ‘idle’ is to call **halt()** defined in **intr.S**. The assembly code of **halt()** only jumps to itself, which means looping forever.

3.2

The **welcome()** function is under **system/welcome.c**.

4 XINU system calls

4.1

The C type of syscall, defined in **include/kernel.h**, is **int**.

4.2

If the parent process has terminated, the process table field **prstate** should be *PR_FREE*.

5 XINU's run-time environment

5.1 Changing byte order using assembly code

The instruction **bswapl** can change byte order of a 32-bit register directly.

5.2 Checking segment boundaries

The addresses of the end of the text, data, and bss segments of the XINU OS are defined by `&etext - 1`, `&edata - 1` and `&ebss - 1` respectively. The **printsegaddress()** function is under `system/printsegaddress.c`.

5.3 Run-time stack: process creation and function call

We will use the inline assembly code to print the top of run-time stack, by copying the value of the register **esp**.

The address and content of the top of the run-time stack inside `main()` will remain the **same** at the time before `myprogA()` is created and after `myprogA()` has been created and resumed. It is because `myprocA()` is a new independent process, and able to run along with the `main()` process. The stack of `myprocA()` will be created in a new free memory area when the process spawns and will **not** change the stack structure of its parent process. Therefore, in these two printing, `main()` process's stack and `main()` function's stack frame are being printed, and the results are same.

When `myprogA()` calls the function `myfuncA()`, things will be different. Before `myfuncA()` is called, `myprocA()` process's stack and `myprocA()` function's stack frame are being printed. Then, `myprogA()` will use its own stack to maintain a function call to `myfuncA()`, pushing the arguments of `myfuncA()` and return address to the stack. After that, `myprogA()` will jump to `myfuncA()`, but the process remains unchanged. In function `myfuncA()`, the stack belongs to the `myprogA()` process. So, after `myfuncA()` is called, `myprocA()` process's stack and `myfuncA()` function's stack frame are being printed.

5.4 Comparing two run-time stacks

As shown in the console.

6 Hijacking a process via stack smashing

The strategy for finding the return address of `sleepms()` function is very simple. In `myfuncA()` process, it is easy to get the process id of `myprogA()` process, i.e. PPID of `myfuncA()` process. Then, we use the PPID to retrieve the stack parameters of `myprogA()` process and the whole stack from `prstkbase` to `prstkptr` can be printed. There are too many addresses at the first glance on the stack information, but the system call `sleepms` has a parameter '3000'. From the system perspective, when a process make a system call or function call, it will push the return address just following the arguments of the function. According to this convention, I find that '0x0000BB8'

(‘3000’ in hex) in the stack and the content under ‘0x00000BB8’ is just the return address of *sleepms* function. In my case, the true return address is ‘0x001034F7’, and the location of the true return address is *proctab[ppid].prstkptr + 120*, since ‘prstkprt’ will be updated to the **esp** register when context switching. We could just replace this content with the *malwareA()*’s address. This offset ‘120’ is fixed because the number of stack allocation for the system call ‘*sleepms()*’ is always ‘120’ bytes. The stack frame of *myprogA()* has been printed in the console.

However, things will not finish. In order to let *myprocA()* process exit normally, which means that its resources can be released and other process with lower priority can continue to run, *myprocA()* process should return to the *INITRET* address after *malwareA()* returns. From the convention, the process will use **ret** instruction to pop the return address from the stack to the **eip** register. Therefore, we may put the *INITRET* address just before the *malwareA()*’s address in the stack. The *INITRET* address is located in *proctab[ppid].prstkbase - 4* defined by *create.c*, so $*(int*)(proctab[ppid].prstkptr + 124) = *(int*)(proctab[ppid].prstkbase - 4)$ would work.

7 Bonus problem

As detailed above, we may rewrite the place just before the *sleepms()* return address in the stack because it is only saving the argument ‘3000’ of *sleepms()*. Thus, we could temporarily save the true return address in that place, i.e.

$$*(int*)(proctab[ppid].prstkptr + 124) = *(int*)(proctab[ppid].prstkptr + 120);$$

Then, move the *malwareB()* address into *proctab[ppid].prstkptr + 120*.

To minimize the disruption on the *myprocA()* process, we cannot modify any register and the content of the running stack after *malwareB()* returns. First, in the normal situation, the address *proctab[ppid].prstkptr + 124* would become useless after the *sleepms()* returns, so saving the true return address there before jumping to *malwareB()* is fine. Second, by looking up the assembly code of *malwareB()*, the **ebx** register which belongs to callee saved has been modified. Third, we cannot just use the true return address saved in *proctab[ppid].prstkptr + 124* to go back to *myprocA()*, because by doing so, the value of the **esp** register will not be equal to the value in the normal situation. Therefore, the only way to make an unusual return is to use inline assembly codes. Before *malwareB()* returns, we should restore the value of **ebx** which has saved in the stack by the function. Then, have the stack frame registers restored normally. Instead of returning now, we must push the true return address from *proctab[ppid].prstkptr + 124* to *proctab[ppid].prstkptr + 120*, in order to imitate the return situation of *sleepms()*. Next, we can do a return by the **ret** instruction. All these return steps have been achieved by inline assembly codes.