

CS50300: Operating Systems

LAB1 ANSWERS

Zichen Wang
wang4113@purdue.edu

September 12, 2018

1 Objectives

2 Readings

3 Running XINU

3.1

An alternative way to ‘idle’ is to call **halt()** defined in **intr.S**. The assembly code of **halt()** only jumps to itself, which means looping forever.

3.2

The **welcome()** function is under **system/welcome.c**.

4 XINU system calls

4.1

The C type of syscall, defined in **include/kernel.h**, is **int**.

4.2

If the parent process has terminated, the process table field **prstate** should be *PR_FREE*.

5 XINU's run-time environment

5.1 Changing byte order using assembly code

The instruction **bswapl** can change byte order of a 32-bit register directly.

5.2 Checking segment boundaries

The addresses of the end of the text, data, and bss segments of the XINU OS are defined by $\&etext - 1$, $\&edata - 1$ and $\&ebss - 1$ respectively. The **printsegaddress()** function is under **system/printsegaddress.c**.

Hints from initialize.c.

5.3 Run-time stack: process creation and function call

We will use the inline assembly code to print the top of run-time stack, by copying the value of the register **esp**.

The address and content of the top of the run-time stack inside **main()** will remain the **same** at the time before **myprogA()** is created and after **myprogA()** has been created and resumed. It is because **myprocA()** is a new independent process, and able to run along with the **main()** process. The stack of **myprocA()** will be created in a new free memory area when the process spawns and will **not** change the stack structure of its parent process. Therefore, in these two printing, **main()** process's stack and **main()** function's stack frame are being printed, and the results are same.

When **myprogA()** calls the function **myfuncA()**, things will be different. Before **myfuncA()** is called, **myprocA()** process's stack and **myprocA()** function's stack frame are being printed. Then, **myprogA()** will use its own stack to maintain a function call to **myfuncA()**, pushing the arguments of **myfuncA()** and return address to the stack. After that, **myprogA()** will jump to **myfuncA()**, but the process remains unchanged. In function **myfuncA()**, the stack belongs to the **myprogA()** process. So, after **myfuncA()** is called, **myprocA()** process's stack and **myfuncA()** function's stack frame are being printed.

5.4 Comparing two run-time stacks

6 Hijacking a process via stack smashing

The strategy for finding the return address of **sleepms()** function is very simple. In **myfuncA()** process, it is easy to get the process id of **myprogA()** process, i.e. PPID of **myfuncA()** process. Then, we use the PPID to retrieve the stack parameters of **myprogA()** process and the whole stack from **prstkbase** to **prstkptr** can be printed. There are too many addresses at the first glance on the stack information, but the system call **sleepms** has a parameter '3000'. From the system perspective, when a process make a system call or function call, it will push the return address just following the arguments of the function. According to this convention, I find that '0x0000BB8' ('3000' in hex) in the stack and the content under '0x0000BB8' is just the return address of **sleepms**

function. In my case, the true return address is '0x001034E4', and the location of the true return address is `proctab[ppid].prstkbase - 52`. We just replace this content by the `malwareA()`'s address.

However, things will not finish. In order to let `myprocA()` process exit normally, which means that its resources can be released and other process with lower priority can continue to run, `myprocA()` process should return to the `INITRET` address after `malwareA()` returns. From the convention, the process will use **ret** instruction to pop the return address from the stack to the **eip** register. Therefore, we may put the `INITRET` address just before the `malwareA()`'s address in the stack. The `INITRET` address is located in `proctab[ppid].prstkbase - 4` defined by `create.c`, so in my case `*(int*)(proctab[ppid].prstkbase - 48) = *(int*)(proctab[ppid].prstkbase - 4)` would work.

7 Bonus problem

As detailed above, we may rewrite the place just before the `sleepms()` return address in the stack because it is only saving the argument '3000' of `sleepms()`. Thus, we could temporarily save the true return address in that place, i.e.

$$*(int*)(proctab[ppid].prstkbase - 48) = *(int*)(proctab[ppid].prstkbase - 52)$$

in my case. Then, move the `malwareB()` address into `proctab[ppid].prstkbase - 52`.

To minimize the influence on the `myprocA()` process, we cannot modify any register and the content of the running stack after `malwareB()` returns. First, in the normal situation, the address `proctab[ppid].prstkbase - 48` would become useless after the `sleepms()` returns, so saving the true return address there before jumping to `malwareB()` is fine. Second, by looking up the assembly code of `malwareB()`, the **ebx** which belongs to callee saved has been modified. Third, we cannot just use the return address saved in `proctab[ppid].prstkbase - 48` to go back to `myprocA()`, because by doing so, the **esp** register will not be equal to the **esp** in the normal situation. Therefore, the only way to make an unusual return is making the use of inline assembly codes. Before `malwareB()` returns, we should restore the value of **ebx** which has saved in the stack. Then, have the stack frame registers restored normally. Instead of returning now, we must push the true return address from `proctab[ppid].prstkbase - 48` to `proctab[ppid].prstkbase - 52`, in order to imitate the return situation of `sleepms()`. Next, we can do a return by **ret** now. All these return steps have been achieved by inline assembly codes.